

AGENTIC AI



YALE SOM

From Chatbots to Agents

- **Chatbot**

- **Input/output:** one-shot prompting process
- **Stateless:** forgets everything when session ends
- **Passive:** only speak when spoke to
- **No tools:** cannot “do” things, only talk about them



- **Agent**

- **Iterative:** loop until goal is met
- **Stateful:** maintain long-term memory
- **Proactive:** monitor environment, act and react
- **Tools:** execute code, write files, call APIs, etc.



How Agents Think

- **Chain of Thought (CoT):** make the agent generate intermediate steps for a query
- **Reason + Act Framework (ReAct):** let the agent think about what action to take and take it


Chain of Thought

- Chain-of-Thought (CoT) was developed in Google in 2022
- Show the AI in its prompt examples of CoT reasoning
- The AI will naturally do its own CoT in its response

Standard Prompting	Chain-of-Thought Prompting
Model Input Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now? A: The answer is 11. Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?	Model Input Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now? A: Roger started with 5 balls. 2 cans of 3 tennis balls each is 6 tennis balls. $5 + 6 = 11$. The answer is 11. Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?
Model Output A: The answer is 27. ❌	Model Output A: The cafeteria had 23 apples originally. They used 20 to make lunch. So they had $23 - 20 = 3$. They bought 6 more apples, so they have $3 + 6 = 9$. The answer is 9. ✅

Reason + Act

- ReAct was developed by Google and Princeton in late 2022
- Combined CoT framework with actions
- Thought->Action->Observe



ReAct Example

(1) Hotspot QA

Question: Aside from the Apple Remote, what other device can control the program Apple Remote was originally designed to interact with?

ReAct Example

(1) Hotspot QA

Question: Aside from the Apple Remote, what other device can control the program Apple Remote was originally designed to interact with?

(1a) Standard

Answer: iPod



ReAct Example

(1) Hotspot QA

Question: Aside from the Apple Remote, what other device can control the program Apple Remote was originally designed to interact with?

(1a) Standard

Answer: iPod



(1b) CoT (Reason Only)

Thought: Let's think step by step. Apple Remote was originally designed to interact with Apple TV. Apple TV can be controlled by iPhone, iPad, and iPod Touch. So the answer is iPhone, iPad, and iPod Touch.

Answer: iPhone, iPad, iPod Touch



ReAct Example

(1) Hotspot QA

Question: Aside from the Apple Remote, what other device can control the program Apple Remote was originally designed to interact with?

(1a) Standard

Answer: iPod



(1b) CoT (Reason Only)

Thought: Let's think step by step. Apple Remote was originally designed to interact with Apple TV. Apple TV can be controlled by iPhone, iPad, and iPod Touch. So the answer is iPhone, iPad, and iPod Touch.

Answer: iPhone, iPad, iPod Touch



(1c) Act-Only

Act 1: Search[Apple Remote]

Obs 1: The Apple Remote is a remote control ...

Act 2: Search[Front Row]

Obs 2: Could not find [Front Row]. Similar: ...

Act 3: Search[Front Row (software)]

Obs 3: Front Row is a discontinued media center software ...

Act 4: Finish[yes]



ReAct Example

(1) Hotspot QA

Question: Aside from the Apple Remote, what other device can control the program Apple Remote was originally designed to interact with?

(1d) ReAct (Reason + Act)

Thought 1: I need to search Apple Remote and find the program it was originally designed to interact with.

Act 1: Search[Apple Remote]

Obs 1: The Apple Remote is a remote control introduced in October 2005 by Apple ... originally designed to control the Front Row media center program ...

Thought 2: Apple Remote was originally designed to control the Front Row media center program. I need to search Front Row next and find what other device can control it.

Act 2: Search[Front Row]

Obs 2: Could not find [Front Row]. Similar: ['Front Row Seat to Earth', 'Front Row Motorsports', 'Front Row (software)', ...]

Thought 3: Front Row is not found. I need to search Front Row (software) .

Act 3: Search[Front Row (software)]

Obs 3: Front Row is a discontinued media center software ...

Thought 4: Front Row (software) is controlled by an Apple Remote or the keyboard function keys. So the answer is keyboard function keys.

Act 4: Finish[keyboard function keys]



Native Tool Calling

- **ReAct lets the AI take actions, but sometimes it would forget the tool format or hallucinate the tool output**
- **Native Tool calling – modern AI models take actions using tools**
- **Train models on code and API documentation**
- **Allows model to pause its text generation and switch to JSON mode instantly**

Building Agents with PydanticAI

- **Pydantic – data validation library for Python**
 - **Makes sure input and output types are correct in your code**
- **PydanticAI – framework built by Pydantic team to make production-grade agents**
 - **Model agnostic – works for Gemini, Claude, GPT**
 - **Dependency injection – safely pass database connections or API keys into agent tools**
 - **Structured results – make sure agents output Python objects**


PydanticAI Code Structure

- **Output format**
- **Agent class**
- **Dependencies**
- **Tools**

PydanticAI Agent Example

- **Price analyst agent**
 - Fetches real-time stock prices
 - Gives you a summary of the current price

What's up
with IBM?



```
✓ ANALYSIS FOR IBM
Price: 237.54 USD
Summary: IBM is currently trading near multi-year highs, driven b
y strong growth in its hybrid cloud and AI segments following sol
id quarterly performance.
```

Output Format

- Define a class to specify how the agent's output is formatted

```
# 2. DEFINE THE OUTPUT SCHEMA
class StockReport(BaseModel):
    symbol: str = Field(description="The stock ticker symbol")
    price: float = Field(description="The current market price")
    currency: str = Field(description="The currency of the price")
    summary: str = Field(description="A brief 1-sentence analyst summary")
```


Dependencies

- Dependencies are how we give the agent state.
- Instead of global variables, we "inject" data safely into the agent's environment.

```
# 2. USER CONTEXT & DEPENDENCIES
@dataclass
class AgentDeps:
    """
    Dependencies injected into the stock analyst agent on each run.
    Provides user preferences and context.
    """
    preferred_currency: str = "USD"
    watchlist: list[str] = field(default_factory=lambda: ["TSLA", "MSFT", "NVDA"])
    user_id: str | None = None
```

Agent Class

- The Agent class manages the conversation state and handles the logic loop between the user and the AI
- Swap between models
- Guarantees agent's answer follows a specific format

```
# 4. THE AGENT
analyst_agent = Agent[AgentDeps, StockReport](
    GEMINI_MODEL,
    output_type=StockReport,
    deps_type=AgentDeps,
    system_prompt=(
        "You are a high-speed financial analyst. "
        "Use the provided tools to find real-time data before answering."
    ),
)
```

Tool Definitions

- Tools are standard Python functions decorated with `@agent.tool` or `@agent.tool_plain`
- The AI uses the Function Name and Docstring to understand when to call it

```
# 5. THE TOOL
@analyst_agent.tool_plain
def get_stock_price(symbol: str) -> str:
    """
    Fetches the current market price for a given stock ticker.
    Args:
        symbol: The stock ticker (e.g., 'AAPL', 'NVDA', 'BTC-USD')
    """
    try:
        ticker = yf.Ticker(symbol)
        # Fetching live data
        price = ticker.fast_info['last_price']
        currency = ticker.fast_info['currency']
        return f"The current price of {symbol} is {price:.2f} {currency}."
    except Exception as e:
        return f"Error fetching price for {symbol}: {str(e)}"
```

Running the Agent

- Define your query

```
query = " ".join(args.query) if args.query else "What is Nvidia's stock  
doing right now?"
```

- Define your agent dependencies

```
deps = AgentDeps(  
    preferred_currency=args.currency,  
    watchlist=args.watchlist if args.watchlist else ["AAPL", "MSFT", "NVDA"],  
    user_id=args.user_id,  
)
```

- Run the agent

```
result = await analyst_agent.run(query, deps=deps)
```

Agent Loop

- **Stock analyst agent does a single pass to complete the task**
 - **Get stock data-> analyze**
- **Advanced agents loop over and over again to complete the given task**
 - **Has access to multiple tools**
 - **Determines in each loop iteration what tool to use**
- **Example: an agent that determines how much of a stock you should buy**

One Agent Loop · Many Tool Choices

Broad goal for **NVDA**: figure out how much to buy. You write **one** `agent.run()` loop — each turn the model picks a tool (or finishes).

● BROAD GOAL

"Figure out how much NVDA to buy."

NVDA

Not a scripted recipe — a goal. The agent decides the steps.

● DEPS (IN THE BACKPACK)

AgentDeps

risk_tolerance
max_position_pct
portfolio_snapshot / user_id

● YOUR CODE

Single call

```
result = await sizer_agent.run(  
    goal, deps=deps  
)
```

PydanticAI runs the loop inside.



Example path (not fixed):

history → news → portfolio → size
Next run might skip news or reorder tools.

● TOOL MENU (PICK ANY, ANY ORDER)

get_price_history

Recent NVDA prices / trend

get_news

Headlines that move the stock

get_portfolio

Cash + current holdings

get_risk_settings

From deps / user risk profile

get_sector_exposure

Already heavy in tech?

Each loop turn: choose one tool or produce the final **PositionAdvice**.

● WHEN THE LOOP STOPS

PositionAdvice

symbol : "NVDA"
action : buy / hold / skip
suggested_dollars
suggested_shares
risk_note
rationale

STRUCTURED EXIT

Pipeline = fixed steps · **Agent loop** = same code, model chooses tools until done

PydanticAI · NVDA Position Sizer

Built-in Agent Tools

- **Models like Gemini come with some useful built-in agent tools you can use**
- **Web search**
- **Code execution**

Web Search Tools

- **What they do**
 - **Grounding:** connects the AI to web search
 - **Freshness:** accesses news, social trends, and real-time data
 - **Verification:** provides citations or links to sources
- **We can give the agent access to web search tools**
 - **Costs money \$\$\$**

Web Search Cost

- Google searches with Gemini cost of \$35 per 1000 searches

Model	Input (per 1M tokens)	Output (per 1M tokens)	Search Grounding Cost
Gemini 3.1 Pro	\$2.00 ($\leq 200k$) / \$4.00 ($> 200k$)	\$12.00 ($\leq 200k$) / \$18.00 ($> 200k$)	\$35 / 1k queries
Gemini 3.5 Flash	\$1.50	\$9.00	\$35 / 1k queries
Gemini 3 Flash	\$0.50	\$3.00	\$35 / 1k queries
Gemini 3.1 Flash-Lite	\$0.25	\$1.50	\$35 / 1k queries

How To Use (API)

- Tool definition:

```
const SEARCH_TOOL = { googleSearch: {} };
```

- Tool call in API:

```
export const streamChat = async function* (history, newMessage, imageParts = [], useCodeExecution = false) {  
  const systemInstruction = await loadSystemPrompt();  
  const tools = useCodeExecution ? [CODE_EXEC_TOOL] : [SEARCH_TOOL];  
  const model = genAI.getGenerativeModel({  
    model: MODEL,  
    tools,  
  });  
};
```

- Tool mention in system prompt

```
You have access to Google Search.  
Use it to verify facts,  
find current events, or provide up-to-date  
technical documentation.
```

Code Execution Tools

- **What it does**
 - **Writes Python code**
 - **Runs code in a secure sandbox**
 - **Use code output to answer your question**
- **This is a basic version of what vibe coders do**

Code Execution Tool Steps

- **Prompt to Gemini:** *I have a CSV with columns timestamp and likes. Can you calculate the correlation between the hour of the day and the number of likes, then tell me the best time to post?*
- **Model's Generated Code (Executable Code)**

```
Python

import pandas as pd
import numpy as np

# 1. Load the data
df = pd.read_csv('engagement_data.csv')

# 2. Extract hour from timestamp
df['hour'] = pd.to_datetime(df['timestamp']).dt.hour

# 3. Calculate Correlation
correlation = df['hour'].corr(df['likes'])

# 4. Find peak performance hour
peak_hour = df.groupby('hour')['likes'].mean().idxmax()

print(f"Correlation: {correlation:.4f}")
print(f"Peak Hour: {peak_hour}")
```

- **Output from sandbox:**

```
Correlation: 0.6421
Peak Hour: 19
```

- **Gemini output:** *Based on the code execution, there is a moderate positive correlation (0.64) between the hour of the day and likes. My analysis shows that your posts perform best at 7:00 PM (19:00). I recommend scheduling your high-value content for that window.*

How to Use (API)

- Tool definition:

```
const CODE_EXEC_TOOL = { codeExecution: {} };  
  
export const CODE_KEYWORDS = /\b(plot|chart|graph|analyz|statistic|regression|correlat
```

- Tool call in API:

```
export const streamChat = async function* (history, newMessage, imageParts = [], useCodeExecution = false) {  
  const systemInstruction = await loadSystemPrompt();  
  const tools = useCodeExecution ? [CODE_EXEC_TOOL] : [SEARCH_TOOL];  
  const model = genAI.getGenerativeModel({  
    model: MODEL,  
    tools,  
  });  
};
```

- Tool mention in system prompt

```
Use code execution (Python) only when the student asks for a  
plot, chart, or something the tools above can't handle.
```

The Power of Code Execution: Iterative Intelligence

- **Think-Execute-Fix Loop:** Unlike a standard chat, the Code Execution tool doesn't just "guess" an answer. It follows a multi-step workflow:
 - **Reasoning:** Gemini identifies that the request requires math or data manipulation.
 - **Generation:** It writes a Python script using libraries like Pandas, NumPy, or Matplotlib.
 - **Self-Correction:** If the code fails (e.g., a missing column or a typo), Gemini sees the error message and automatically rewrites the code. It can try up to 5 times before giving up.
- **Built-in Data Science Toolkit:** the sandbox comes pre-loaded with the industry-standard libraries:
 - **Pandas:** For cleaning messy social media exports.
 - **Matplotlib:** For generating charts (scatter plots, histograms) on the fly.
 - **NumPy:** For high-performance statistical calculations.

Be Careful



Summer Yue  @summeryue0 · Feb 22



Nothing humbles you like telling your OpenClaw “confirm before acting” and watching it speedrun deleting your inbox. I couldn’t stop it from my phone. I had to RUN to my Mac mini like I was defusing a bomb.


Be Careful



Summer Yue  @summeryue0 · Feb 22



Nothing humbles you like telling your OpenClaw “confirm before acting” and watching it speedrun deleting your inbox. I couldn’t stop it from my phone. I had to RUN to my Mac mini like I was defusing a bomb.





Be Careful




Summer Yue @summeryue0 · Feb 22

Nothing humbles you like telling your OpenClaw “confirm before acting” and watching it speedrun deleting your inbox. I couldn’t stop it from my phone. I had to RUN to my Mac mini like I was defusing a bomb.



Be Careful



A screenshot of a Twitter profile page for Summer Yue. The header features a large background image of a person walking a dog on a beach at sunset. On the left is a circular profile picture of Summer Yue, a woman with long dark hair wearing a grey beanie and a red flower. To the right of the profile picture are three icons: a three-dot menu, an envelope, and a 'Follow' button. Below the profile picture, the name 'Summer Yue' is displayed with a blue verification checkmark, followed by the handle '@summeryue0'. The bio text is: 'Safety and alignment at Meta Superintelligence. Prev: VP of Research at Scale AI, research at Google DeepMind / Brain (Gemini, LaMDA, RL / TFAgents, AlphaChip)'. The location 'San Francisco, CA' and 'Joined August 2014' are shown below the bio. At the bottom, it says '387 Following' and '16.1K Followers'.

Summer Yue ✓
@summeryue0

Safety and alignment at Meta Superintelligence. Prev: VP of Research at Scale AI, research at Google DeepMind / Brain (Gemini, LaMDA, RL / TFAgents, AlphaChip).

San Francisco, CA Joined August 2014 >

387 Following 16.1K Followers

Be Extra Careful


OpenAI admits its models hacked Hugging Face on their own

They escaped an isolated environment for testing and infiltrated Hugging Face without human input.


by Mariella Moon • July 21, 2026 11:47 pm EST




OpenAI was doing **red-team** testing on GPT-5.6 Sol in a sandbox



GPT realizes it needs some data outside the sandbox, so it escapes from the sandbox



GPT realizes the data it needs is on HuggingFace, so it breaks into HuggingFace



Building Agents with Agents

- PydanticAI has a lot of syntax
- Vibe coding agents (i.e. Cursor) know the syntax or can look it up with a web search tool
- So we can build agents with agents

